
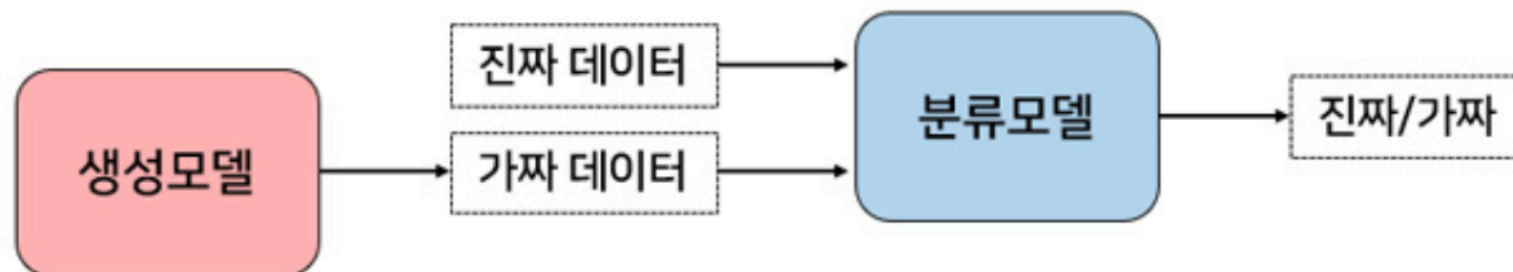
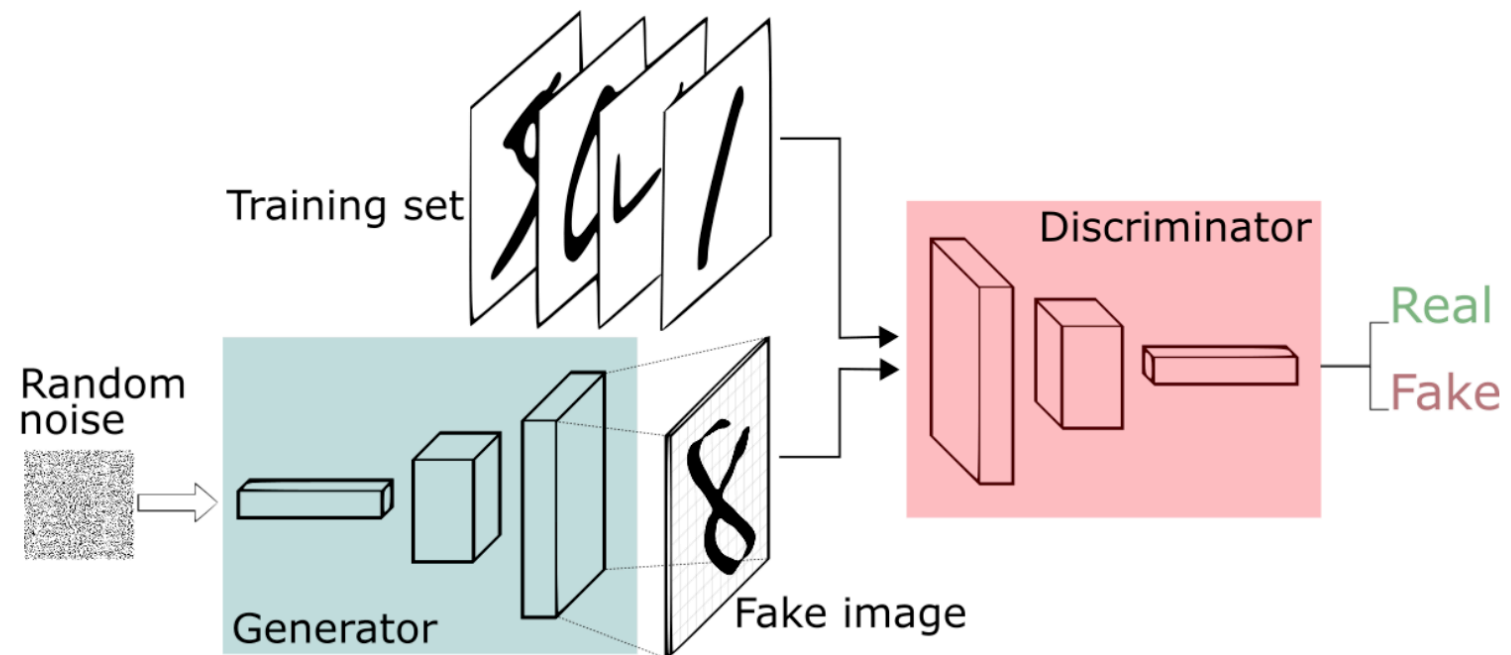
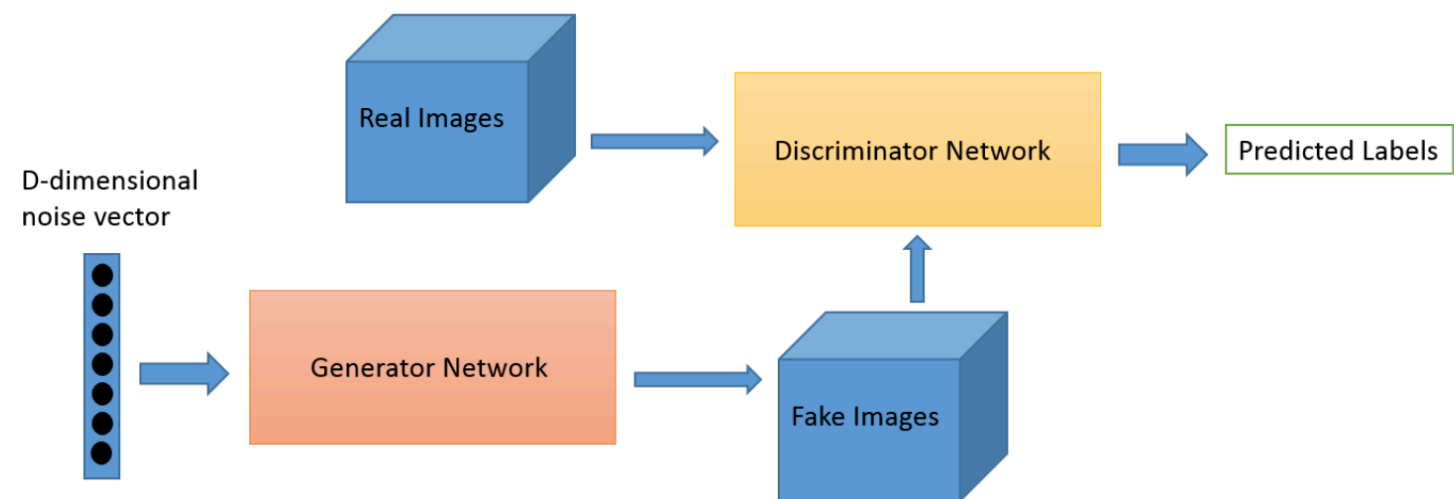



DCGAN

UNSUPERVISED REPRESENTATION
LEARNING WITH DEEP CONVOLUTIONAL
GENERATIVE ADVERSARIAL NETWORKS









Architecture guidelines for stable Deep Convolutional GANs

- Replace any pooling layers with strided convolutions (discriminator) and fractional-strided convolutions (generator).
- Use batchnorm in both the generator and the discriminator.
- Remove fully connected hidden layers for deeper architectures.
- Use ReLU activation in generator for all layers except for the output, which uses Tanh.
- Use LeakyReLU activation in the discriminator for all layers.

No pre-processing was applied to training images besides scaling to the range of the tanh activation function $[-1, 1]$.

All models were trained with mini-batch stochastic gradient descent (SGD) with a mini-batch size of 128.

All weights were initialized from a zero-centered Normal distribution with standard deviation 0.02.

In the LeakyReLU, the slope of the leak was set to 0.2 in all models.

While previous GAN work has used momentum to accelerate training, we used the Adam optimizer (Kingma & Ba, 2014) with tuned hyperparameters.

We found the suggested learning rate of 0.001, to be too high, using 0.0002 instead.

Additionally, we found leaving the momentum term β_1 at the suggested value of 0.9 resulted in training oscillation and instability while reducing it to 0.5 helped stabilize training.

| Layer (type) | Output Shape | Param # |
|---|---------------------|----------|
| dense_1 (Dense) | (None, 16384) | 2113536 |
| batch_normalization_1 (Batch Normalization) | (None, 16384) | 65536 |
| re_lu_1 (ReLU) | (None, 16384) | 0 |
| reshape_1 (Reshape) | (None, 4, 4, 1024) | 0 |
| dropout_1 (Dropout) | (None, 4, 4, 1024) | 0 |
| up_sampling2d_1 (UpSampling2D) | (None, 8, 8, 1024) | 0 |
| conv2d_transpose_1 (Conv2DTranspose) | (None, 8, 8, 512) | 13107712 |
| batch_normalization_2 (Batch Normalization) | (None, 8, 8, 512) | 2048 |
| activation_1 (Activation) | (None, 8, 8, 512) | 0 |
| dropout_2 (Dropout) | (None, 8, 8, 512) | 0 |
| up_sampling2d_2 (UpSampling2D) | (None, 16, 16, 512) | 0 |
| conv2d_transpose_2 (Conv2DTranspose) | (None, 16, 16, 256) | 3277056 |
| batch_normalization_3 (Batch Normalization) | (None, 16, 16, 256) | 1024 |
| activation_2 (Activation) | (None, 16, 16, 256) | 0 |
| dropout_3 (Dropout) | (None, 16, 16, 256) | 0 |
| up_sampling2d_3 (UpSampling2D) | (None, 32, 32, 256) | 0 |
| conv2d_transpose_3 (Conv2DTranspose) | (None, 32, 32, 128) | 819328 |
| batch_normalization_4 (Batch Normalization) | (None, 32, 32, 128) | 512 |
| re_lu_2 (ReLU) | (None, 32, 32, 128) | 0 |
| dropout_4 (Dropout) | (None, 32, 32, 128) | 0 |
| up_sampling2d_4 (UpSampling2D) | (None, 64, 64, 128) | 0 |
| conv2d_transpose_4 (Conv2DTranspose) | (None, 64, 64, 3) | 9603 |
| activation_3 (Activation) | (None, 64, 64, 3) | 0 |

| Layer (type) | Output Shape | Param # |
|---|---------------------|---------|
| conv2d_1 (Conv2D) | (None, 64, 64, 512) | 14336 |
| batch_normalization_5 (Batch Normalization) | (None, 64, 64, 512) | 2048 |
| leaky_re_lu_1 (LeakyReLU) | (None, 64, 64, 512) | 0 |
| dropout_5 (Dropout) | (None, 64, 64, 512) | 0 |
| conv2d_2 (Conv2D) | (None, 32, 32, 512) | 2359808 |
| batch_normalization_6 (Batch Normalization) | (None, 32, 32, 512) | 2048 |
| leaky_re_lu_2 (LeakyReLU) | (None, 32, 32, 512) | 0 |
| dropout_6 (Dropout) | (None, 32, 32, 512) | 0 |
| conv2d_3 (Conv2D) | (None, 16, 16, 128) | 589952 |
| batch_normalization_7 (Batch Normalization) | (None, 16, 16, 128) | 512 |
| leaky_re_lu_3 (LeakyReLU) | (None, 16, 16, 128) | 0 |
| dropout_7 (Dropout) | (None, 16, 16, 128) | 0 |
| conv2d_4 (Conv2D) | (None, 8, 8, 128) | 147584 |
| batch_normalization_8 (Batch Normalization) | (None, 8, 8, 128) | 512 |
| leaky_re_lu_4 (LeakyReLU) | (None, 8, 8, 128) | 0 |
| dropout_8 (Dropout) | (None, 8, 8, 128) | 0 |
| flatten_1 (Flatten) | (None, 8192) | 0 |
| dense_2 (Dense) | (None, 1) | 8193 |
| activation_4 (Activation) | (None, 1) | 0 |
| Total params: 3,124,993 | | |
| Trainable params: 3,122,433 | | |
| Non-trainable params: 2,560 | | |

| Layer (type) | Output Shape | Param # |
|---------------------------|-------------------|----------|
| sequential_1 (Sequential) | (None, 64, 64, 3) | 19396355 |
| sequential_2 (Sequential) | (None, 1) | 3124993 |

Total params: 22,521,348
 Trainable params: 19,361,795
 Non-trainable params: 3,159,553

```

def load_data(self, dataset_name, flag=1):
    if flag==0:
        TRAIN_DIR = os.path.join("./dataset", dataset_name)
        print(TRAIN_DIR)
        training_data = []
        for img in os.listdir(TRAIN_DIR):
            # print(img)
            path = os.path.join(TRAIN_DIR, img)
            # print(path)
            img = cv2.imread(path, cv2.IMREAD_COLOR)
            # print(img)
            img = cv2.resize(img, (self.img_size, self.img_size))
            training_data.append([np.array(img).astype('float32')])
        shuffle(training_data)
        x_train = np.vstack(training_data) / 255.0
        np.save(str(dataset_name) + '_train_data.npy', x_train)
        print(x_train.shape)
    else:
        x_train=np.load(str(dataset_name)+'_train_data.npy')
  
```



```
def build_model(self):  
  
    Image_Data_Class = ImageData(self.img_size, self.c_dim)  
    self.data = Image_Data_Class.load_data(dataset_name=self.dataset_name)  
    self.input_shape=(self.img_size,self.img_size,self.c_dim)  
  
    self.g_model=self.generator()  
    self.g_model.summary()  
    self.d_model=self.discriminator()  
    self.d_model.summary()  
  
    self.g_optimizer = Adam(lr=self.g_learning_rate)  
    self.g_model.compile(self.g_optimizer,loss="binary_crossentropy")  
  
    self.g_d_model = Sequential()  
    self.g_d_model.add(self.g_model)  
    self.d_model.trainable = False  
    self.g_d_model.add(self.d_model)  
    self.g_d_model.summary()  
  
    self.g_d_optimizer = Adam(lr=self.g_learning_rate)  
    self.g_d_model.compile(self.g_d_optimizer,loss="binary_crossentropy")  
  
    self.d_model.trainable = True    d_learning_rate  
    self.d_optimizer = Adam(lr=self.d_learning_rate)  
    self.d_model.compile(self.d_optimizer,loss="binary_crossentropy")
```

```

def train(self):
    x_train = self.data
    print(x_train.shape)
    BATCH_SIZE=self.batch_size

    # Some parameters.

    for epoch in range(self.epoch):
        print("Epoch is", epoch)
        print("Number of batches", int(x_train.shape[0]/BATCH_SIZE))
        for index in range(int(x_train.shape[0]/BATCH_SIZE)):
            noise = np.random.uniform(-1.0, 1.0, size=(BATCH_SIZE, self.z_dim))
            image_batch = x_train[index*BATCH_SIZE:(index+1)*BATCH_SIZE]
            generated_images = self.g_model.predict(noise, verbose=0)
            X = np.concatenate((image_batch, generated_images))
            y = np.ones([2*BATCH_SIZE, 1])
            y[BATCH_SIZE:,:] = 0
            self.d_model.trainable=True
            self.d_loss = self.d_model.train_on_batch(X, y)
            print("batch %d d_loss : %f" % (index, self.d_loss))

            y = np.ones([BATCH_SIZE, 1])
            noise = np.random.uniform(-1.0, 1.0, (BATCH_SIZE, self.z_dim))
            self.d_model.trainable = False
            self.g_loss = self.g_model.train_on_batch(noise, y)
            print("batch %d g_loss : %f" % (index, self.g_loss))
            if epoch % 10 == 0:
                self.g_model.save_weights('generator', True)
                self.d_model.save_weights('discriminator', True)

    self.phase='test'

```

```
def test(self):
    g = self.generator()
    g.compile(loss='binary_crossentropy', optimizer="SGD")
    g.load_weights('generator')
    for i in range(20):
        noise = np.random.uniform(-1.0, 1.0, (self.batch_size, self.z_dim))
        generated_images = g.predict(noise, verbose=1)
        image = image*255.0
        Image.fromarray(image.astype(np.uint8)).save("generated_image_"+str(i) + ".png")
```



sequential_2 (Sequential) (None, 1) 3124993

Total params: 6,382,788
Trainable params: 3,224,579
Non-trainable params: 3,158,209

(7864, 64, 64, 3)
('Epoch is', 0)
('Number of batches', 78)
batch 0 d_loss : 7.971192
batch 0 g_loss : 0.000000
batch 1 d_loss : 7.971192
batch 1 g_loss : 0.000000

□